

Dataset Description

A person makes a doctor appointment, receives all the instructions and no-show. Who to blame? This dataset collects information from 100k medical appointments in Brazil and is focused on the question of whether or not patients show up for their appointment. A number of characteristics about the patient are included in each row. ____

Columns Description

1. `PatientId` : Identification of a patient.
 2. `AppointmentID` : Identification of each appointment.
 3. `Gender` : Male or Female.
 4. `AppointmentDay` : The day of the actual appointment, when they have to visit the doctor.
 5. `ScheduledDay` : The day someone called or registered the appointment, this is before appointment of course.
 6. `Age` : How old is the patient.
 7. `Neighbourhood` : Where the appointment takes place.
 8. `Scholarship` : True or False, indicates whether or not the patient is enrolled in Brazilian welfare program Bolsa Família.
 9. `Hipertension` : True or False.
 10. `Diabetes` : True or False.
 11. `Alcoholism` : True or False.
 12. `Handcap` : True or False.
 13. `SMS_received` : 1 or more messages sent to the patient.
 14. `No-show` : True (if the patient did not show up), or False (if the patient did show up).
-

EDA Questions

Q1: How often do men go to hospitals compared to women? Which of them is more likely to show up?

Q2: Does receiving an SMS as a reminder affect whether or not a patient may show up? is it correlated with number of days before the appointment?

Q3: Does having a scholarship affect showing up on a hospital appointment? What are the age groups affected by this?

Q4: Does having certain diseases affect whether or not a patient may show up to their appointment? is it affected by gender? ____

Environment set-up

```
In [1]: # import required libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib.ticker as ticker
```

Data Wrangling

We'd load our desired data from the flat csv file `noshowappointments - kaggle2-may-2016.csv` to a dataframe using `pandas`, and display its first 5 records. here, we want to check for:

- Missingness in our dataframe.
- Inconsistent data types.
- NaNs.
- Duplicated rows.
- columns to be dropped or re-parsed.

```
In [2]: #Load Data
df = pd.read_csv(r'D:\Projects\My portfolio - Completed projects\Python\2024 10 Medical Appointment Dataset An

#check top rows
df.head(5)
```

	PatientId	AppointmentID	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood	Scholarship	Hipertension	Diabet
0	2.987250e+13	5642903	F	2016-04-29T18:38:08Z	2016-04-29T00:00:00Z	62	JARDIM DA PENHA	0		1
1	5.589978e+14	5642503	M	2016-04-29T16:08:27Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	0		0
2	4.262962e+12	5642549	F	2016-04-29T16:19:04Z	2016-04-29T00:00:00Z	62	MATA DA PRAIA	0		0
3	8.679512e+11	5642828	F	2016-04-29T17:29:31Z	2016-04-29T00:00:00Z	8	PONTAL DE CAMBURI	0		0
4	8.841186e+12	5642494	F	2016-04-29T16:07:23Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	0		1

```
In [3]: # display the number of rows and columns in the dataset
df.shape
```

Out[3]: (110527, 14)

```
In [4]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 110527 entries, 0 to 110526
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   PatientId             110527 non-null float64
1   AppointmentID          110527 non-null int64
2   Gender                 110527 non-null object
3   ScheduledDay           110527 non-null object
4   AppointmentDay         110527 non-null object
5   Age                   110527 non-null int64
6   Neighbourhood          110527 non-null object
7   Scholarship            110527 non-null int64
8   Hipertension           110527 non-null int64
9   Diabetes               110527 non-null int64
10  Alcoholism             110527 non-null int64
11  Handcap                110527 non-null int64
12  SMS_received           110527 non-null int64
13  No-show                110527 non-null object
dtypes: float64(1), int64(8), object(5)
memory usage: 11.8+ MB
```

We can notice there are no NaNs at all in our data:

- `ScheduledDay` and `AppointmentDay` needs to be casted to date data type.
- we may append a new column for days until appointment.
- `Gender` needs to be converted into a category type
- `Scholarship` `Hipertension` `Diabetes` `Alcoholism` `Handcap` better be boolean data type.
- `No-show` needs to be parsed and casted to boolean too.

```
In [5]: # checking for duplicates
df.duplicated().sum()
```

Out[5]: 0

our dataset has no duplicated rows either.

```
In [6]: # exploring the unique values of each column
df.nunique()
```

```
Out[6]: PatientId      62299
AppointmentID    110527
Gender              2
ScheduledDay      103549
AppointmentDay     27
Age               104
Neighbourhood      81
Scholarship         2
Hipertension        2
Diabetes            2
Alcoholism          2
Handcap             5
SMS_received        2
No-show            2
dtype: int64
```

```
In [7]: # exploring handicap values
df.Handcap.value_counts()
```

```
Out[7]: Handcap
0      108286
1       2042
2        183
3         13
4          3
Name: count, dtype: int64
```

```
In [8]: # exploring Age values
df.Age.value_counts()
```

```
Out[8]: Age
0      3539
1      2273
52     1746
49     1652
53     1651
...
115      5
100      4
102      2
99       1
-1       1
Name: count, Length: 104, dtype: int64
```

```
In [9]: # exploring age column distribution
df.Age.describe()
```

```
Out[9]: count      110527.000000
mean         37.088874
std          23.110205
min          -1.000000
25%          18.000000
50%          37.000000
75%          55.000000
max          115.000000
Name: Age, dtype: float64
```

- `Handcap` column has inconsistent unique values, we'd be only interested in rows with `0` or `1` values.
- `Age` column has inconsistent unique values, we'd better handle it.
- `SMS_received` would be casted to boolean data type.

Exploration Summary

1. our dataset consists of 110527 rows with 14 columns, and has no NaNs nor duplicated values.
2. `PatientId` and `AppointmentId` columns wouldn't be helpful during analysis.
3. `ScheduledDay` and `AppointmentDay` needs to be casted to date data type.
4. we may append a new column for days until appointment.
5. `Gender` needs to be casted into a category type
6. `Scholarship`, `Hipertension`, `Diabetes`, `Alcoholism` and `SMS_recieved` better be boolean data type.
7. `No-show` column needs to be parsed and casted to boolean type.
8. `Handcap` column needs to be cleaned to have only `0` and `1` values.
9. `Age` column has inconsistent unique values that needs to be handled.

Data Cleaning

in this section, we'd perform some operations on our dataset based on the previous findings to make our analysis more accurate and clear.

Dropping `PatientId` and `AppointmentId` columns

```
In [10]: df.drop(['PatientId', 'AppointmentID'], axis = 1, inplace = True)

df.columns
```

```
Out[10]: Index(['Gender', 'ScheduledDay', 'AppointmentDay', 'Age', 'Neighbourhood',
               'Scholarship', 'Hipertension', 'Diabetes', 'Alcoholism', 'Handcap',
               'SMS_received', 'No-show'],
              dtype='object')
```

Handling `date` data type

```
In [11]: df.AppointmentDay.unique
```

```
Out[11]: <bound method Series.unique of 0          2016-04-29T00:00:00Z
1          2016-04-29T00:00:00Z
2          2016-04-29T00:00:00Z
3          2016-04-29T00:00:00Z
4          2016-04-29T00:00:00Z
...
110522     2016-06-07T00:00:00Z
110523     2016-06-07T00:00:00Z
110524     2016-06-07T00:00:00Z
110525     2016-06-07T00:00:00Z
110526     2016-06-07T00:00:00Z
Name: AppointmentDay, Length: 110527, dtype: object>
```

it looks like all hours are set to 00:00:00, so we would want to extract the year, month and day data

```
In [12]: # extracting only day, month and year values
df['ScheduledDay'] = df['ScheduledDay'].astype(str).str[:10]
df['AppointmentDay'] = df['AppointmentDay'].astype(str).str[:10]

# changing data type
df['ScheduledDay'] = pd.to_datetime(df['ScheduledDay'])
df['AppointmentDay'] = pd.to_datetime(df['AppointmentDay'])

# check changes
print(df[['AppointmentDay', 'ScheduledDay']].dtypes)
```

```
AppointmentDay    datetime64[ns]
ScheduledDay      datetime64[ns]
dtype: object
```

```
In [13]: df.head()
```

```
Out[13]:
```

	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood	Scholarship	Hipertension	Diabetes	Alcoholism	Handcap	SMS
0	F	2016-04-29	2016-04-29	62	JARDIM DA PENHA	0	1	0	0	0	
1	M	2016-04-29	2016-04-29	56	JARDIM DA PENHA	0	0	0	0	0	
2	F	2016-04-29	2016-04-29	62	MATA DA PRAIA	0	0	0	0	0	
3	F	2016-04-29	2016-04-29	8	PONTAL DE CAMBURI	0	0	0	0	0	
4	F	2016-04-29	2016-04-29	56	JARDIM DA PENHA	0	1	1	0	0	

Now, we'd move into appending a new column that holds number of days to the appointment.

```
In [14]: # making new due days column
df['DueDays'] = df['AppointmentDay'] - df['ScheduledDay']

# converting data type
df['DueDays'] = df['DueDays'].dt.days
```

```
In [15]: df.DueDays
```

```
Out[15]: 0          0
1          0
2          0
3          0
4          0
...
110522     35
110523     35
110524     41
110525     41
110526     41
Name: DueDays, Length: 110527, dtype: int64
```

```
In [16]: # viewing summary statistics
df['DueDays'].describe()
```

```
Out[16]: count      110527.000000
         mean        10.183702
         std         15.254996
         min         -6.000000
         25%         0.000000
         50%         4.000000
         75%        15.000000
         max        179.000000
         Name: DueDays, dtype: float64
```

We seem to have some negative values here, we'll drop them.

```
In [17]: df[df['DueDays'] < 0]
```

```
Out[17]:
```

	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood	Scholarship	Hipertension	Diabetes	Alcoholism	Handcap
27033	M	2016-05-10	2016-05-09	38	RESISTÊNCIA	0	0	0	0	1
55226	F	2016-05-18	2016-05-17	19	SANTO ANTÔNIO	0	0	0	0	1
64175	F	2016-05-05	2016-05-04	22	CONSOLAÇÃO	0	0	0	0	0
71533	F	2016-05-11	2016-05-05	81	SANTO ANTÔNIO	0	0	0	0	0
72362	M	2016-05-04	2016-05-03	7	TABUAZEIRO	0	0	0	0	0

```
In [18]: # dropping negative values
         df.drop(df[df['DueDays'] < 0].index, inplace = True)
         df['DueDays'].describe()
```

```
Out[18]: count      110522.000000
         mean        10.184253
         std         15.255115
         min         0.000000
         25%         0.000000
         50%         4.000000
         75%        15.000000
         max        179.000000
         Name: DueDays, dtype: float64
```

Converting Gender to categorical variables

```
In [19]: # converting gender column
         df['Gender'] = df['Gender'].astype('category')

         df['Gender'].dtypes
```

```
Out[19]: CategoricalDtype(categories=['F', 'M'], ordered=False, categories_dtype=object)
```

Parsing and casting No-show column

```
In [20]: # mapping alues to be more familiar
         df.loc[df['No-show'] == 'Yes', 'No-show'] = 0
         df.loc[df['No-show'] == 'No', 'No-show'] = 1

         # casting dt type and confirming changes
         df['No-show'] = df['No-show'].astype(bool)
         df['No-show'].dtypes
```

```
Out[20]: dtype('bool')
```

Converting Scholarship, Hipertension, Diabetes, Alcoholism, Handcap and SMS_received to boolean data type

```
In [21]: # converting columns to bool
         columns = ['Scholarship', 'Hipertension', 'Diabetes', 'Alcoholism', 'SMS_received']
         df[columns] = df[columns].astype(bool)
         df[columns].dtypes
```

```
Out[21]: Scholarship      bool
         Hipertension      bool
         Diabetes          bool
         Alcoholism        bool
         SMS_received       bool
         dtype: object
```

Cleaning Handcap column

```
In [22]: # check rows with values of handicap > 1
```

```
df[df['Handcap'] > 1]
```

Out[22]:

	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood	Scholarship	Hipertension	Diabetes	Alcoholism	Handcap
946	M	2016-04-14	2016-04-29	94	BELA VISTA	False	True	True	False	2
1665	M	2016-03-30	2016-04-29	64	SANTA MARTHA	False	True	False	True	2
1666	M	2016-03-30	2016-04-29	64	SANTA MARTHA	False	True	False	True	2
2071	M	2016-04-29	2016-04-29	64	SANTA MARTHA	False	True	False	True	2
2091	F	2016-04-29	2016-04-29	11	ANDORINHAS	False	False	False	False	2
...
108376	F	2016-06-01	2016-06-07	44	ROMÃO	False	True	True	False	2
109484	M	2016-05-31	2016-06-02	64	DA PENHA	False	True	True	False	2
109733	F	2016-06-03	2016-06-07	34	JUCUTUQUARA	False	False	False	False	2
109975	M	2016-06-02	2016-06-06	39	PRAIA DO SUÁ	True	False	False	False	2
110107	F	2016-06-02	2016-06-06	44	RESISTÊNCIA	False	False	False	False	2

199 rows × 13 columns

◀												▶
---	--	--	--	--	--	--	--	--	--	--	--	---

We have 199 rows with inconsistant values, we'd replace them with 1 to treat them as beeing handicaped.

```
In [23]: # filling the bigger values with 1
df.loc[df['Handcap'].isin([2, 3, 4]), 'Handcap'] = 1

# casting type
df['Handcap'] = df['Handcap'].astype(bool)
```

```
In [24]: df['Handcap'].unique()
```

Out[24]: array([False, True])

Cleaning Age column

```
In [25]: #checking negative values
df[df['Age'] < 0]
```

Out[25]:

	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood	Scholarship	Hipertension	Diabetes	Alcoholism	Handcap
99832	F	2016-06-06	2016-06-06	-1	ROMÃO	False	False	False	False	False

◀												▶
---	--	--	--	--	--	--	--	--	--	--	--	---

We have one value with negative age, so we will drop it.

```
In [26]: # dropping row with negative age
df.drop(df[df['Age'] < 0].index, inplace = True)

#checking negative values
df[df['Age'] < 0]
```

Out[26]:

	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood	Scholarship	Hipertension	Diabetes	Alcoholism	Handcap	SMS_
--	--------	--------------	----------------	-----	---------------	-------------	--------------	----------	------------	---------	------

◀												▶
---	--	--	--	--	--	--	--	--	--	--	--	---

Checking data again

```
In [27]: df.head()
```

Out[27]:

	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood	Scholarship	Hipertension	Diabetes	Alcoholism	Handcap	SMS
0	F	2016-04-29	2016-04-29	62	JARDIM DA PENHA	False	True	False	False	False	
1	M	2016-04-29	2016-04-29	56	JARDIM DA PENHA	False	False	False	False	False	
2	F	2016-04-29	2016-04-29	62	MATA DA PRAIA	False	False	False	False	False	
3	F	2016-04-29	2016-04-29	8	PONTAL DE CAMBURI	False	False	False	False	False	
4	F	2016-04-29	2016-04-29	56	JARDIM DA PENHA	False	True	True	False	False	

In [28]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 110521 entries, 0 to 110526
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                 110521 non-null  category
1   ScheduledDay           110521 non-null  datetime64[ns]
2   AppointmentDay         110521 non-null  datetime64[ns]
3   Age                   110521 non-null  int64
4   Neighbourhood          110521 non-null  object
5   Scholarship            110521 non-null  bool
6   Hipertension           110521 non-null  bool
7   Diabetes               110521 non-null  bool
8   Alcoholism             110521 non-null  bool
9   Handcap                110521 non-null  bool
10  SMS_received           110521 non-null  bool
11  No-show                110521 non-null  bool
12  DueDays                110521 non-null  int64
dtypes: bool(7), category(1), datetime64[ns](2), int64(2), object(1)
memory usage: 5.9+ MB
```

We ended up with a dataframe of 110521 rows and 13 columns, and everything looks tidy and clean. We'd proceed in visualizing it to extract meaningful insights from it. _____

Data Visualization and EDA

Now that our data is clean, we'd perform some EDA on it in order to extract useful insights from it.

In [29]: `# setting seaborn configurations`
`sns.set_style("dark")`
`palette_options = ['bright']`

1. How often do men go to hospitals compared to women? Which of them is more likely to show up?

In [30]: `df['Gender'].value_counts()`


Out[30]:

```
Gender
F    71836
M    38685
Name: count, dtype: int64
```


In [31]: `# Display the distribution of the gender column`
`plt.figure(figsize=(8, 7))`
`ax = sns.countplot(x='Gender', hue='Gender', data=df, palette='bright', legend=False)`
`plt.title("Gender Distribution")`
`ax.set_xlabel('Gender')`
`ax.set_ylabel('Count')`

`# Adding numerical labels on each bar and formatting`
`for p in ax.patches:`
 `count = int(p.get_height())`
 `ax.annotate(f'{count}',`
 `(p.get_x() + p.get_width() / 2., count),`
 `ha='center', va='bottom',`
 `xytext=(0, 5),`
 `textcoords='offset points')`

`plt.show()`



```
In [32]: # Display gender distribution against age in our dataset
plt.figure(figsize=(8, 4))
ax = sns.boxplot(x='Gender', y='Age', data=df)
plt.title("Gender Distribution against Age")
plt.show()
```



- we can notice that nearly half of our dataset consists of women with wider age distribution and some outliers, all of which achieving a rate higher than men.

```
In [33]: no_show_counts = df['No-show'].value_counts()
no_show_percentages = df['No-show'].value_counts(normalize=True) * 100

# Combine counts and percentages into a DataFrame
no_show_summary = pd.DataFrame({
    'Count': no_show_counts,
    'Percentage': no_show_percentages
})

no_show_summary
```


Out[33]:

	Count	Percentage
--	-------	------------

No-show

True	88207	79.810172
------	-------	-----------

False	22314	20.189828
-------	-------	-----------

- it is obvious that 79.8% of our patients did show up on their appointments and only 20.1% of them did not.

let's dive deeper to see if this is related to gender.

```
In [34]: gender_no_show_table = pd.crosstab(df['Gender'], df['No-show'])

# Rename columns
gender_no_show_table.columns = ['Show', 'No-show']

gender_no_show_table
```

Out[34]:

	Show	No-show
--	------	---------

Gender

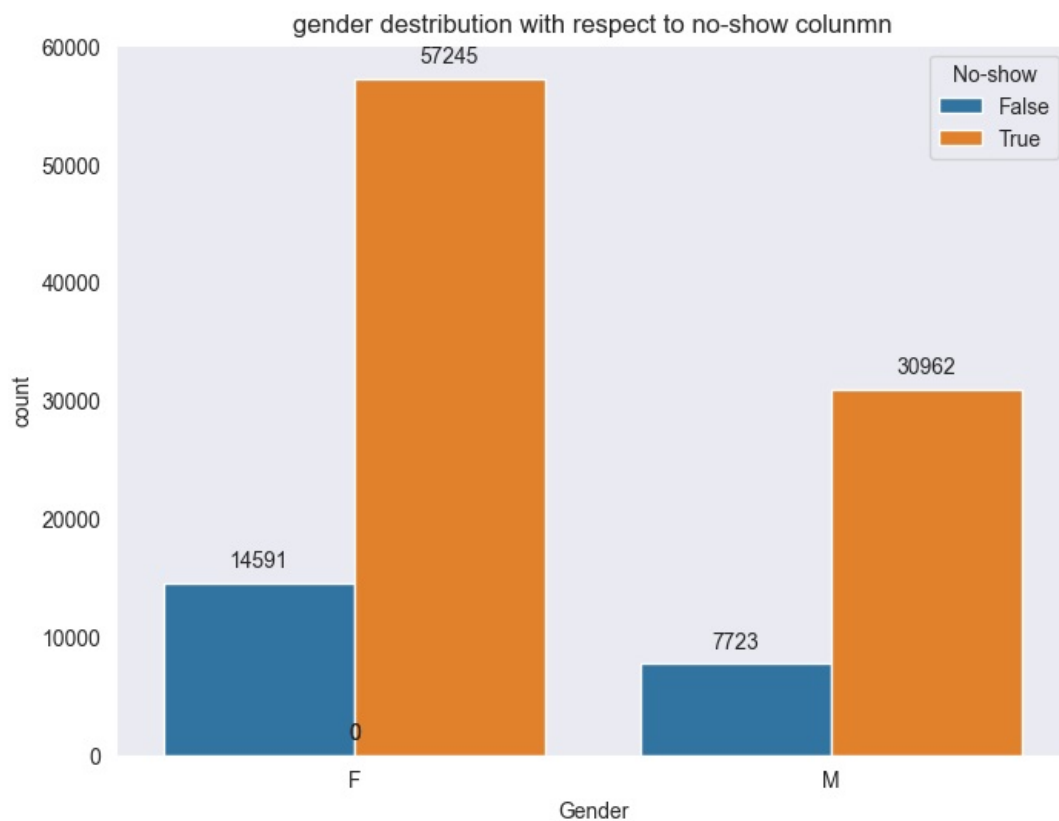
F	14591	57245
---	-------	-------

M	7723	30962
---	------	-------

```
In [35]: # Display gender distribution with respect to the no-show columnn
plt.figure(figsize=(8, 6))
ax = sns.countplot(x = 'Gender', data = df, hue = 'No-show')
plt.title('gender distribution with respect to no-show columnn')

# Adding numerical labels on each bar and formatting
for p in ax.patches:
    count = int(p.get_height())
    ax.annotate(f'{count}',
                (p.get_x() + p.get_width() / 2., count),
                ha='center', va='bottom',
                xytext=(0, 5),
                textcoords='offset points')

plt.show()
```



- from the above chart, we can come up with a conclusion that women do show up on their appointments more often than men do, but this may be affected by the percentage of women on this dataset.

before the appointment?

```
In [36]: # recieving SMS distribution
SMS_received_counts = df['SMS_received'].value_counts()
SMS_received_percentages = df['SMS_received'].value_counts(normalize=True)* 100

# Combine counts and percentages into a DataFrame
SMS_received_summary = pd.DataFrame({
    'Count': SMS_received_counts,
    'Percentage': SMS_received_percentages
})

SMS_received_summary
```

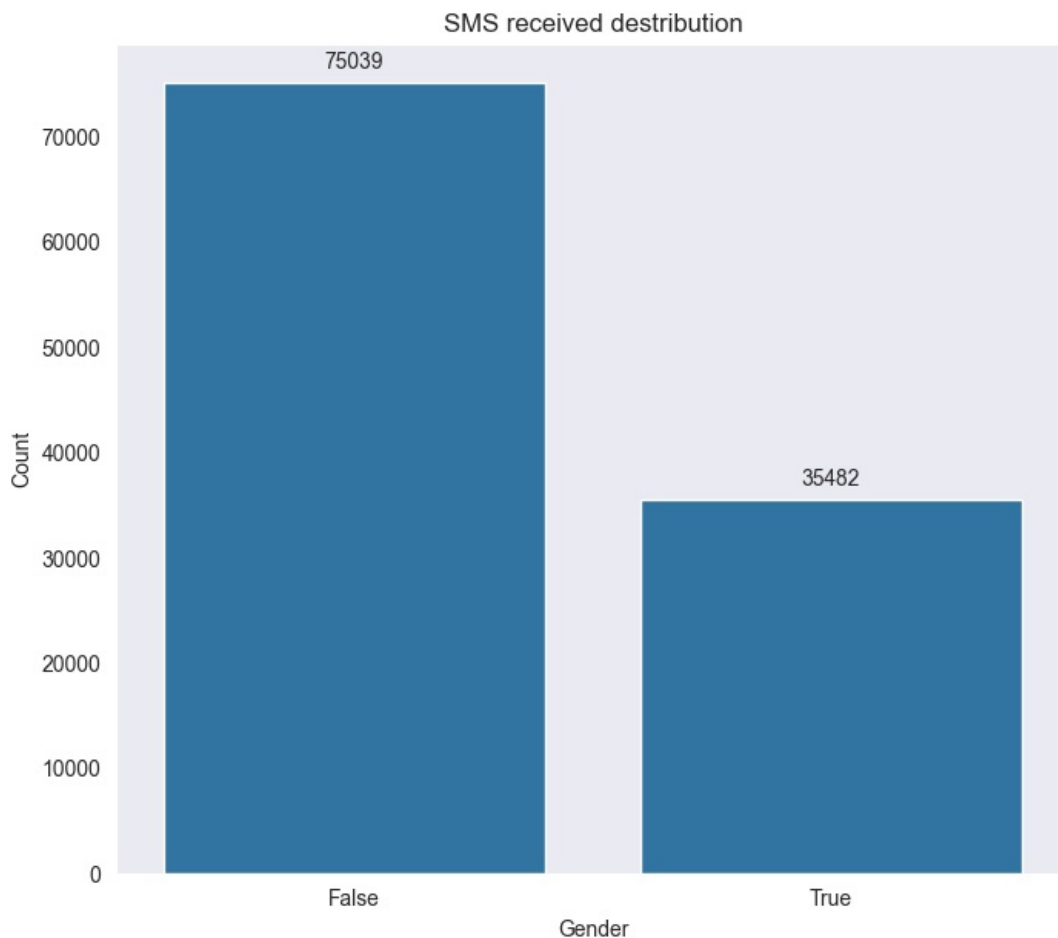
```
Out[36]:
```

	Count	Percentage
SMS_received		
False	75039	67.895694
True	35482	32.104306

```
In [37]: # Display recieving SMS distribution
plt.figure(figsize=(8, 7))
ax = sns.countplot(x='SMS_received', data=df)
plt.title("SMS received distribution")
ax.set_xlabel('Gender')
ax.set_ylabel('Count')

# Adding numerical labels on each bar and formatting
for p in ax.patches:
    count = int(p.get_height())
    ax.annotate(f'{count}',
                (p.get_x() + p.get_width() / 2., count),
                ha='center', va='bottom',
                xytext=(0, 5),
                textcoords='offset points')

plt.show()
```



- we can see that 67.8% of our patients did not recive any SMS reminder of their appointments, is this may be affecting their showin up?

```
In [38]: SMS_received_no_show_table = pd.crosstab(df['SMS_received'], df['No-show'])
```

```
SMS_received_no_show_table.columns = ['SMS_received', 'No-show']
```

```
SMS_received_no_show_table
```

```
Out[38]:
```

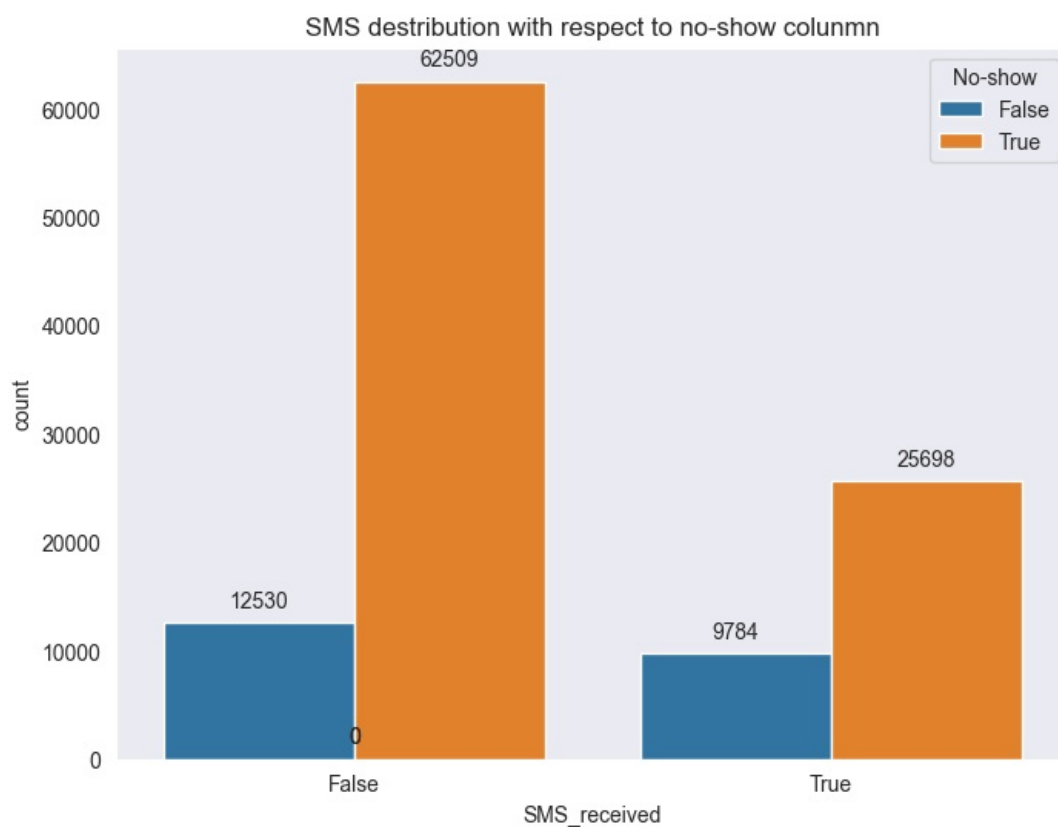
	SMS_received	No-show
SMS_received		
False	12530	62509
True	9784	25698

```
In [39]:
```

```
# Display gender distribution with respect to the no-show columnn
plt.figure(figsize=(8, 6))
ax = sns.countplot(x = 'SMS_received', data = df, hue = 'No-show')
plt.title('SMS distribution with respect to no-show columnn')

# Adding numerical labels on each bar and formatting
for p in ax.patches:
    count = int(p.get_height())
    ax.annotate(f'{count}',
                (p.get_x() + p.get_width() / 2., count),
                ha='center', va='bottom',
                xytext=(0, 5),
                textcoords='offset points')


plt.show()
```



- we can see that our previous deduction was not quiet correct, as the vast majority of our patients did not recieve any SMS reminder and yet they showed up on their appointments.

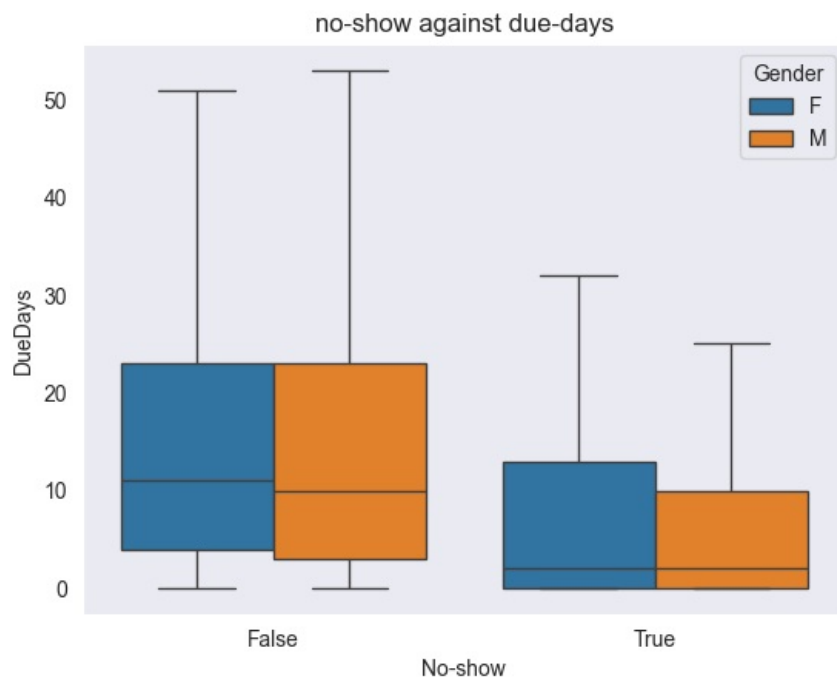
```
In [40]:
```

```
# viewing the correlation between no-show and due-days without outliers
sns.boxplot(x = 'No-show', y = 'DueDays', data = df, showliers = False)
plt.title('no-show against due-days')
plt.show()
```



- it is clear that there is a correlation between number of due days and whether a patient shows up or not.
- patient with appointments from 0 to 30 days tend to show up more regularly, while patients with higher number of days tend to not show up.

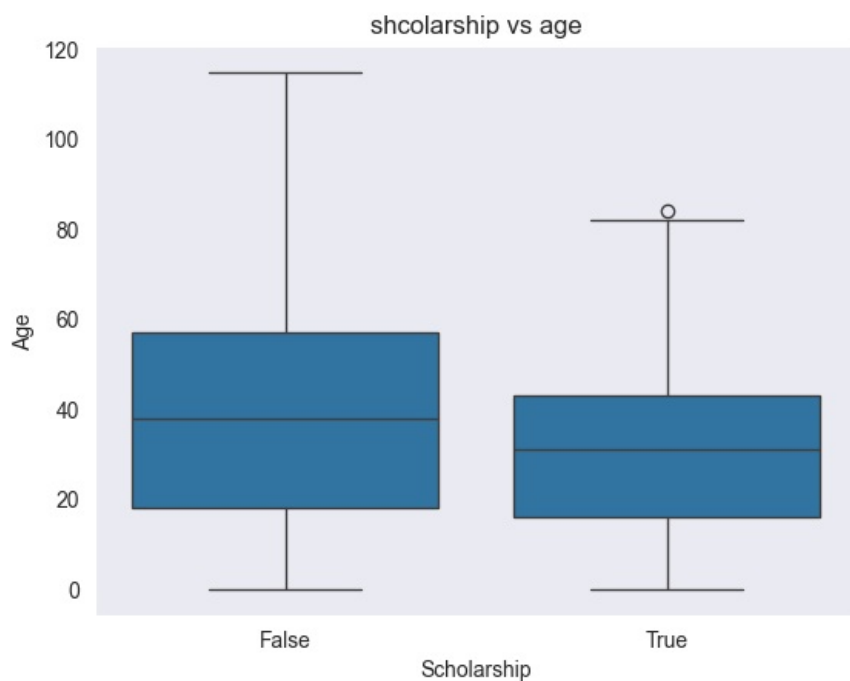
```
In [41]: # viewing the correlation between no-show and due-days without outliers with respect to gender
sns.boxplot(x = 'No-show', y = 'DueDays', data = df, hue = 'Gender', showfliers = False)
plt.title('no-show against due-days')
plt.show()
```



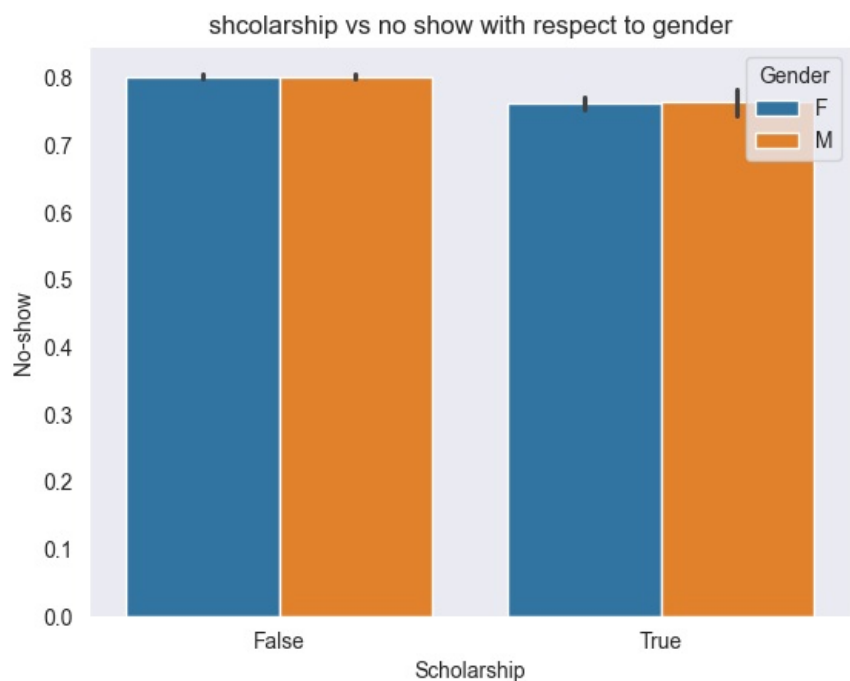
- gender does not affect number of due days and showing up at an appointment that much.

3. Does having a scholarship affects showing up on a hospital appointment? What are the age groups affected by this?

```
In [42]: # Display scholarship against age
sns.boxplot(x = 'Scholarship', y = 'Age', data = df)
plt.title('scholarship vs age')
plt.show()
```




```
In [43]: # plotting having a scholarship against no show with respect to gender
ax = sns.barplot(x = 'Scholarship', y = 'No-show', hue = 'Gender', data = df)
plt.title('shcolarship vs no show with respect to gender')
plt.show()
```



```
In [44]: # plotting age destribution
ax = df['Age'].hist(bins=10, edgecolor='black', figsize=(8, 6))
plt.title('Age Distribution')
plt.xlabel('Age')
plt.ylabel('Count')

for p in ax.patches:
    plt.text(p.get_x() + p.get_width() / 2,
             p.get_height(),
             f'{int(p.get_height())}',
             ha='center',
             va='bottom',)

plt.show()
```

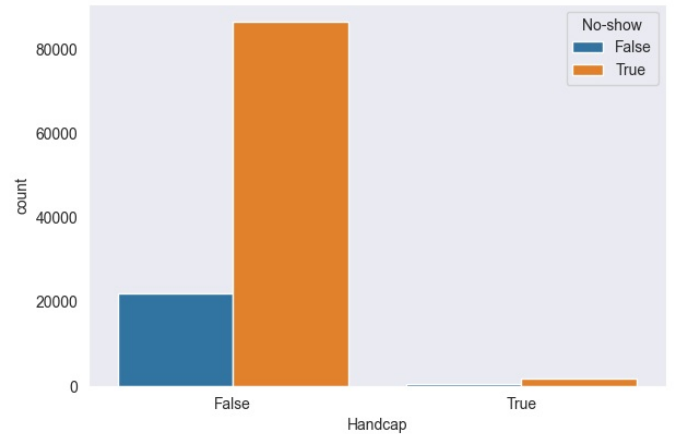
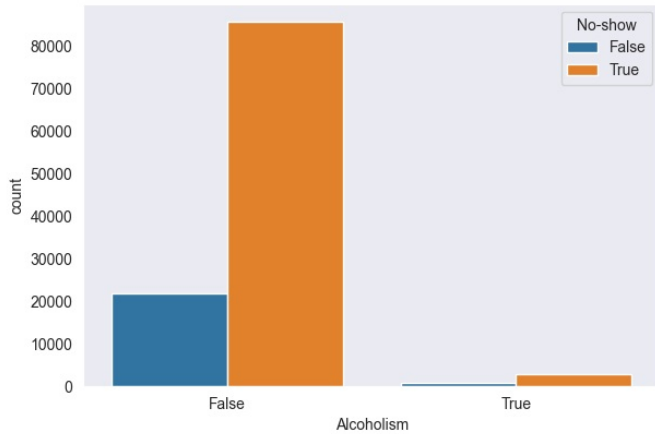
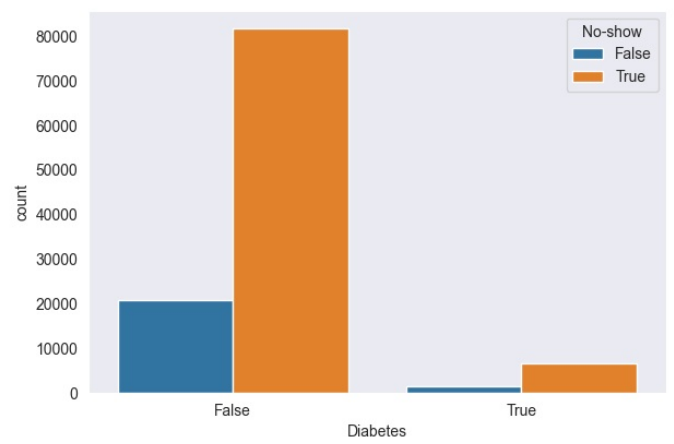
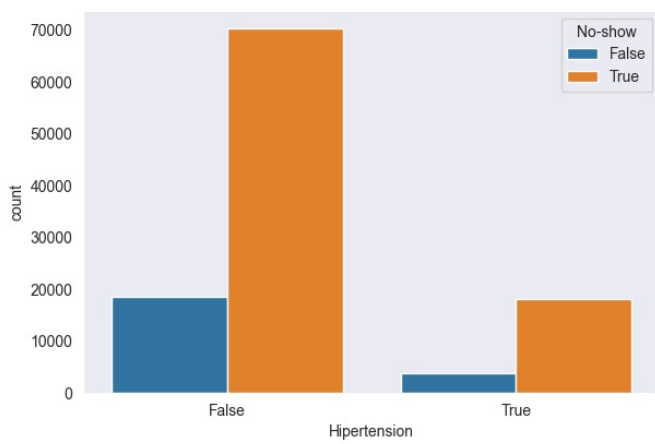


- we can see that having a scholarship doesn't affect showing up to a doctor appointment that much and that huge age group is enrolled to that scholarship and also enrol their babies on.

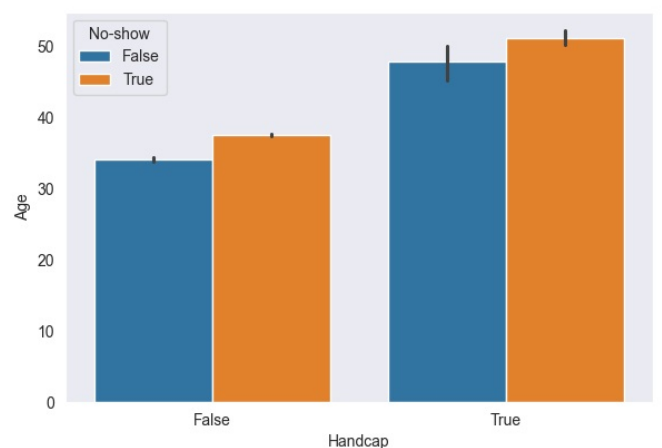
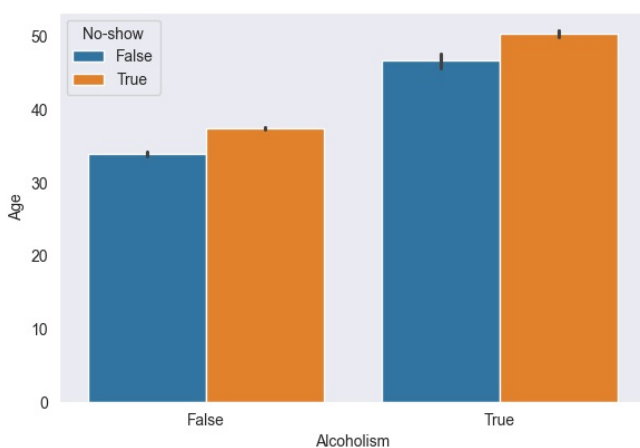

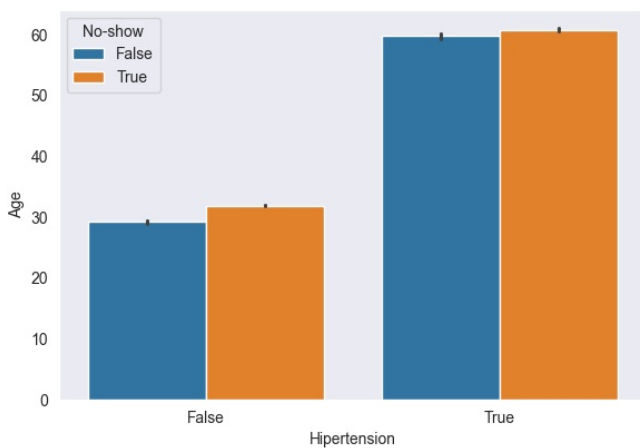
4. Does having certain diseases affects whather or not a patient may show up to their appointment? is it affected by gender?

In [45]: # plotting diseases against no show

```
plt.figure(figsize=(15,10))
plt.subplot(2,2,1)
sns.countplot(x = 'Hipertension', data = df, hue= 'No-show')
plt.subplot(2,2,2)
sns.countplot(x = 'Diabetes', data = df, hue= 'No-show')
plt.subplot(2,2,3)
sns.countplot(x = 'Alcoholism', data = df, hue= 'No-show')
plt.subplot(2,2,4)
sns.countplot(x = 'Handcap', data = df, hue= 'No-show')
plt.show()
```



```
In [46]: # plotting diseases against no show with respect to age
plt.figure(figsize=(15,10))
plt.subplot(2,2,1)
sns.barplot(x = 'Hipertension', y = 'Age', data = df, hue= 'No-show')
plt.subplot(2,2,2)
sns.barplot(x = 'Diabetes', y = 'Age', data = df, hue= 'No-show')
plt.subplot(2,2,3)
sns.barplot(x = 'Alcoholism', y = 'Age', data = df, hue= 'No-show')
plt.subplot(2,2,4)
sns.barplot(x = 'Handcap', y = 'Age', data = df, hue= 'No-show')
plt.show()
```



- from the previous set of plots, we can conclude that the vast majority of our dataset does not have chronic diseases, yet, they are existed in so many young people.
 - having a chronic disease may affect your showing up at a hospital's appointment.
-

Conclusion

Q1: How often do men go to hospitals compared to women? Which of them is more likely to show up?

- Nearly half of our dataset consists of women with wider age distribution and some outliers, all of which achieve a rate higher than men.
 - It is obvious that 79.8% of our patients did show up on their appointments and only 20.1% of them did not.
 - Women do show up on their appointments more often than men do, but this may be affected by the percentage of women on this dataset.
-

Q2: Does receiving an SMS as a reminder affect whether or not a patient may show up? is it correlated with number of days before the appointment?

- 67.8% of our patients did not receive any SMS reminder of their appointments, yet they showed up on their appointments.
 - It is clear that there is a positive correlation between number of due days and whether a patient shows up or not.
 - Patient with appointments from 0 to 30 days tend to show up more regularly, while patients with higher number of days tend to not show up.
 - gender does not affect number of due days and showing up at an appointment that much.
-

Q3: Does having a scholarship affect showing up on a hospital appointment? What are the age groups affected by this?

- Having a scholarship does not affect showing up to a doctor appointment that much.
 - Huge age group is enrolled to that scholarship and also enrol their babies on.
-

Q4: Does having certain diseases affect whether or not a patient may show up to their appointment? is it affected by gender?

- We can conclude that the vast majority of our dataset does not have chronic diseases, yet, they are existed in so many young people.
- Having a chronic disease may affect your showing up at a hospital's appointment.

"This project was entirely developed by Bassam El-Shoraa".